

# Supervised learning

- Data Preparation and Splitting
  - Train, validation, test approach
  - K-fold Cross validation approach
- Regression
  - Bias-variance tradeoff
  - Relationship with model complexity
  - Underfitting and Overfitting
  - Linear regression
  - Evaluation metrics

# Splitting the data



Original

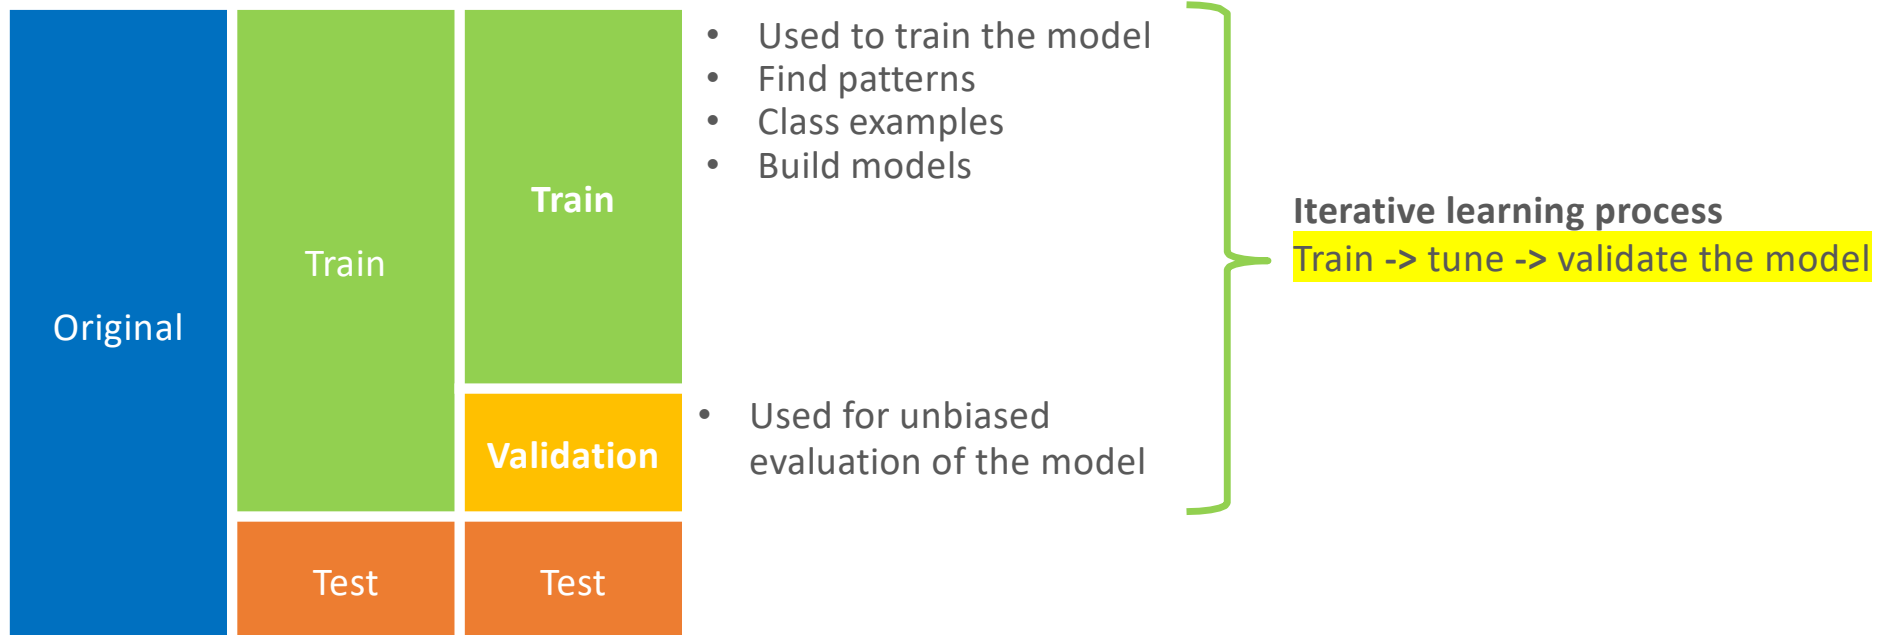
- Ensure that the instances are representative and relevant to the problem
- Ensure collaboration in data preprocessing
- *It's a good practice to shuffle the data before the split to avoid bias in the resulting sets*

# Train, test, and validation split approach


In most supervised machine learning tasks, best practice recommends to split your data into three independent sets:

- Training set
- Validation set
- Testing set (It must not be touched until the end)

# Train, test, and validation split approach



# Train, test, and validation split approach




- It is not available to the model for learning
- It is only used to ensure that the model generalizes well on new “unseen” data
- It is used for final evaluation of the model

# K-Fold Cross-Validation approach

- This strategy iteratively use different portions (folds) of our data to train and validate our model
- Useful when data is limited and we want more robust estimates of model performance

# K-Fold Cross-Validation approach (e.g., $k=4$ )




The final **test** is evaluated by average performance on the test set metrics across all the folds

1. Split the **train** into  $k$  independent folds
2. Repeat the following  $k$  times:
  1. Set aside  $k^{th}$  fold (hold-out set) of the data for **validation**
  2. Train the model on the remaining  $k - 1$  folds
  3. Test the model on the **validation**
3. At the end, average or combine the model performance metrics

\*Test data remains unchanged

<https://erickedu85.github.io>

# Process preview




Choice a ML algorithm

- Suitable for desired label
- Suitable for learning type
- Suitable for data size
- Suitable given available info
- Suitable given practical issues
  - Production requirements
  - Engineering effort
  - Expertise required



# Process preview



Choice a ML algorithm

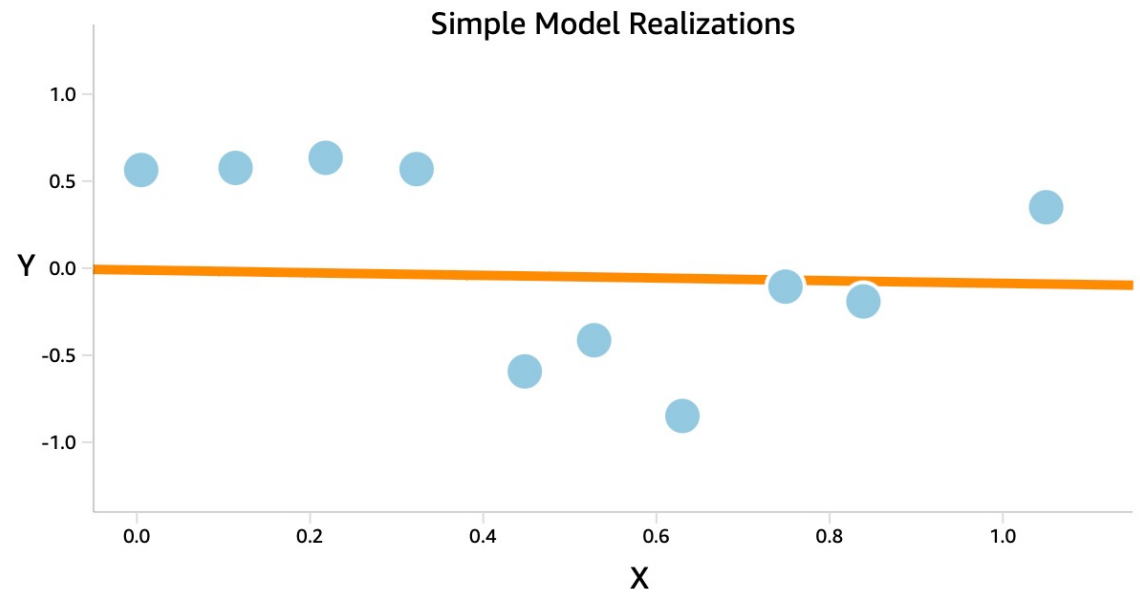
- Suitable for desired label
- Suitable for learning type
- Suitable for data size
- Suitable given available info
- Suitable given practical issues
  - Production requirements
  - Engineering effort
  - Expertise required

# The bias-variance tradeoff

- Simple model

# A simple model


No matter how we try to fit this line, *the straight line does not have the flexibility to accurately capture the relationship of the data*



The model is too simple to capture the structure of the data

# A simple model

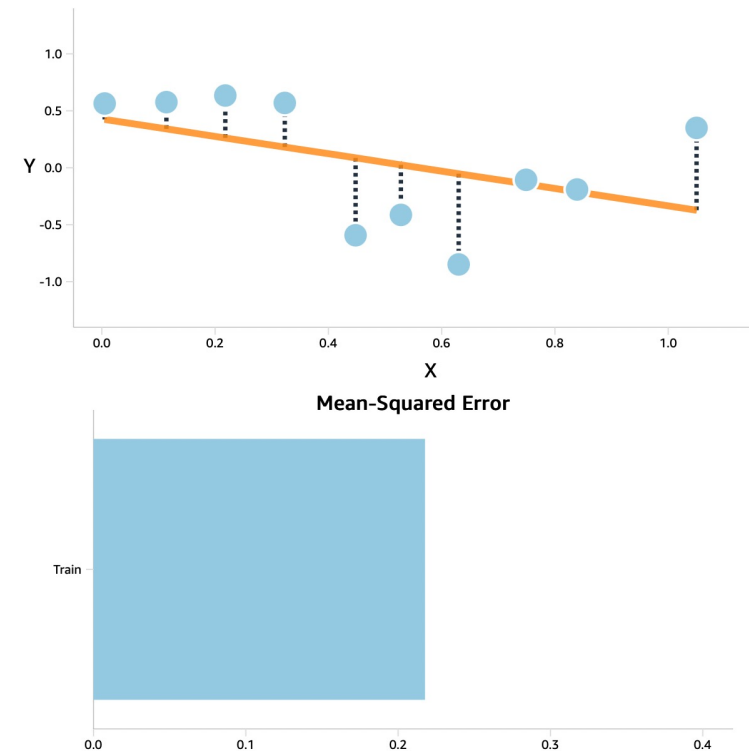
No matter how we try to fit this line, *the straight line does not have the flexibility to accurately capture the relationship of the data*



The model is too simple to capture the structure of the data


# A simple model

- No matter how we try to fit this line, *the straight line does not have the flexibility to accurately capture the relationship of the data*
- Error in the **training data** is not zero



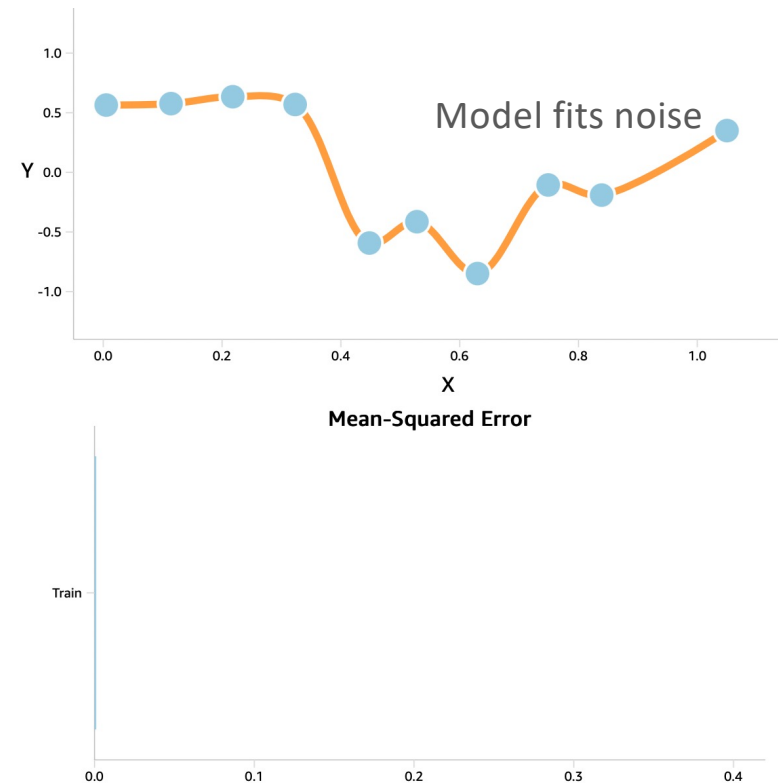
# Underfitting: low complexity model

- High error on **train data** and **test data**
- The model is so simple that it fails to adequately capture the relationships in the data
- The high error on test is a direct result of the lack of complexity of the model



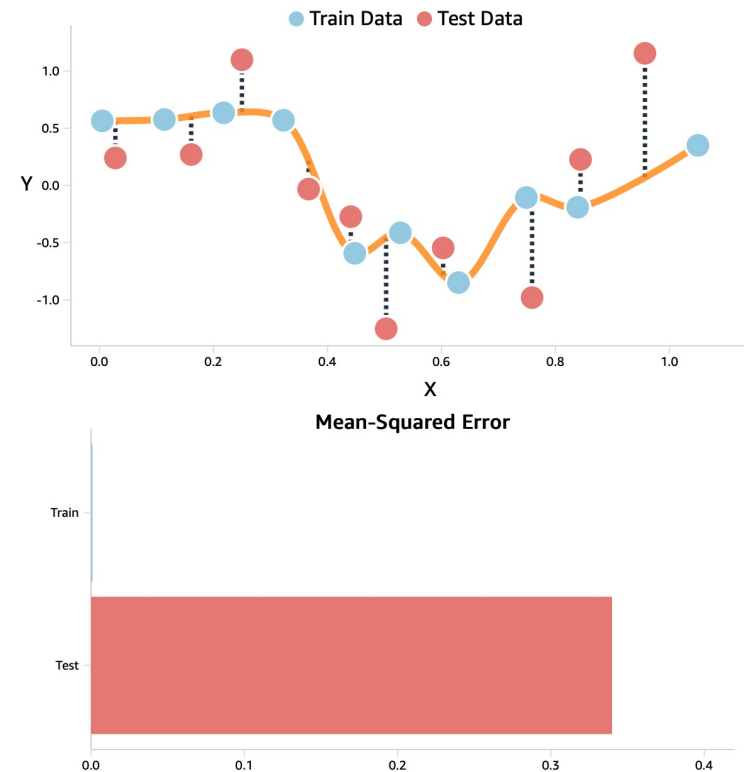
# A complex model

- A more complex model predicts every point in the training data perfectly
- Super flexible model
- Error in the **train data** is zero



# Overfitting: high complex model

- **Low error** on **train data** and **high** on **test data**
- The model is so specific to the data on which it was training that it is no longer applicable to different datasets
- The model is not generalizable to other datasets (e.g., **test data**) ⚠





# Test error decomposition

Mean Squared Error (MSE) can be decomposed into three components:


- $Bias^2$  : Error due to wrong assumptions
- *Variance*: Error due to sensitivity to data
- *Irreducible error*: Error due to noise in data

$$Error = Bias^2 + Variance + Noise$$

$$Error = (E[\hat{f}(x)] - f(x))^2 + E[(\hat{f}(x) - E[\hat{f}(x)])^2] + Noise$$

We can make use of the relationship between both bias and variance to obtain better predictions.

Test Error Decomposition



# Bias and variance

## Bias

- It measures how far the **average** model prediction is from the **actual** outcome
- High bias usually means the model is too **simple** to capture the underlying patterns (underfitting)

$$\text{Bias}(x) = E[\hat{f}(x)] - f(x)$$

$\hat{f}(x)$ : the prediction made by a trained model

$E[\hat{f}(x)]$ : the expected prediction from many models trained on different random training datasets

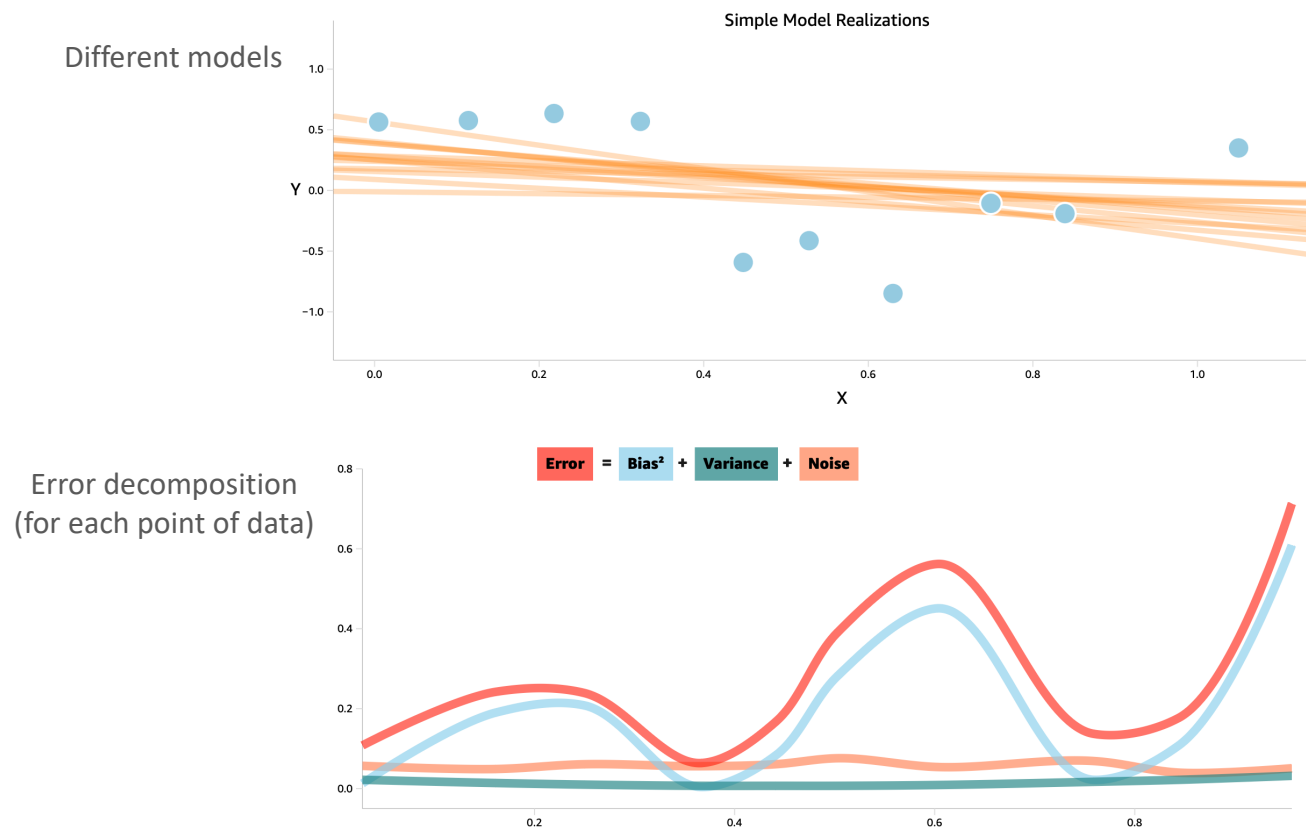
$f(x)$ : the true value (true function or ground truth)

## Variance

- It measures how much the model's prediction **change** when trained on different samples of the data
- High variance often means the model is too **complex** and sensitive to noise (overfitting)


$$\text{Variance}(x) = E[(\hat{f}(x) - E[\hat{f}(x)])^2]$$

# Bias and variance



For underfit models (low-complexity), the majority of the error comes from **bias**

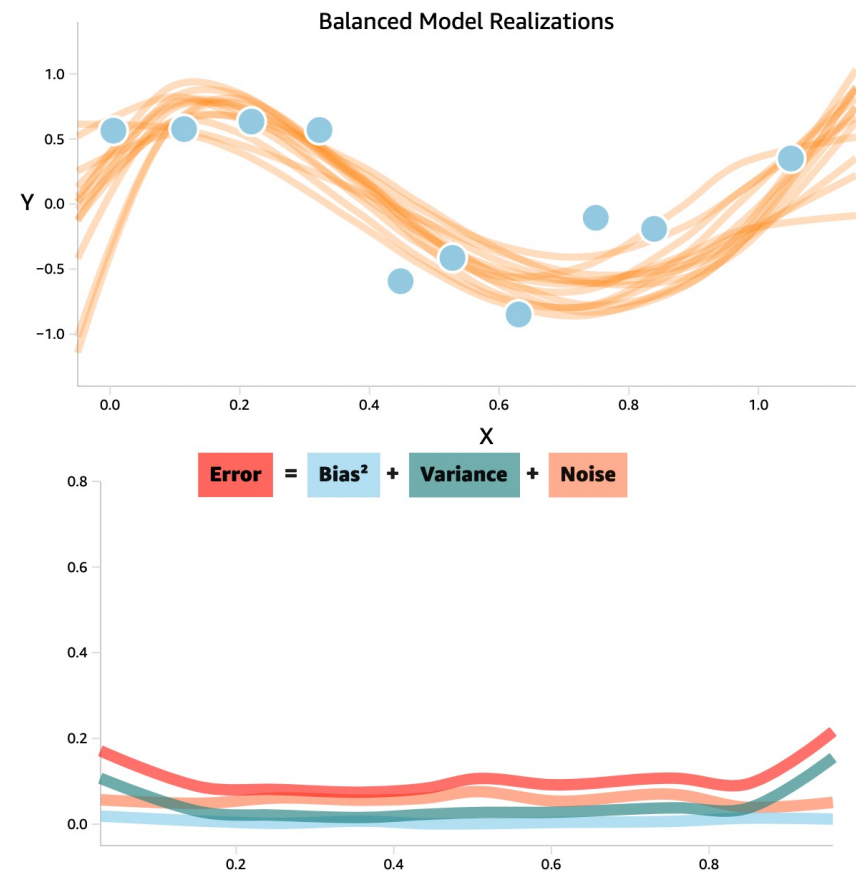
# Bias and variance



Overfit models (high-complexity) show a lot more error from **variance** than from bias. It's easy to imagine that any unseen data points will be predicted with high error

# Find a balance

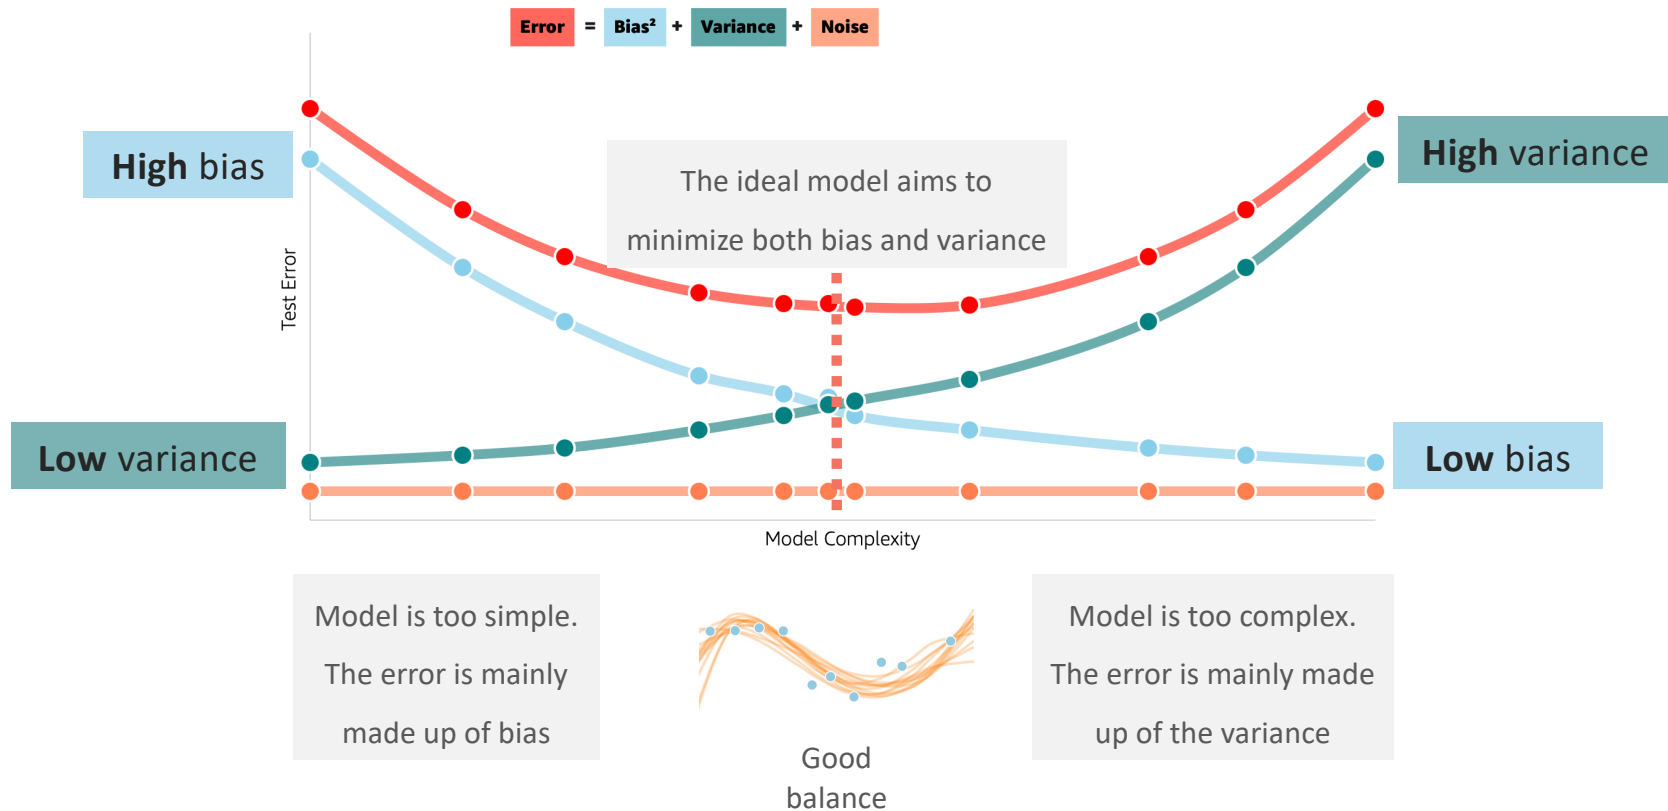
- Find a balance between **underfitting** (model too simple to learn meaningful patterns) and **overfitting** (model too complex to generalize to unseen data).
- By allowing a bit more variance (e.g., increasing model complexity) to reduce bias, we can aim for a model with optimal generalization performance.



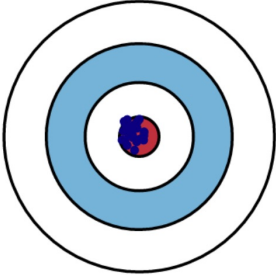
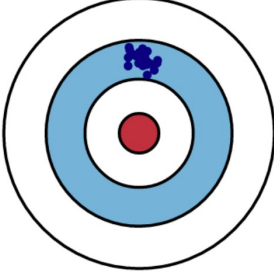
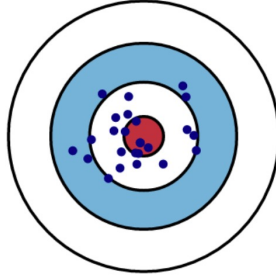
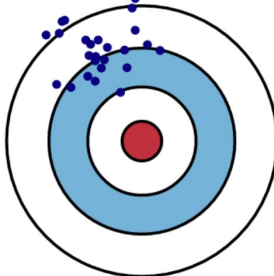
# Bias-Variance Tradeoff

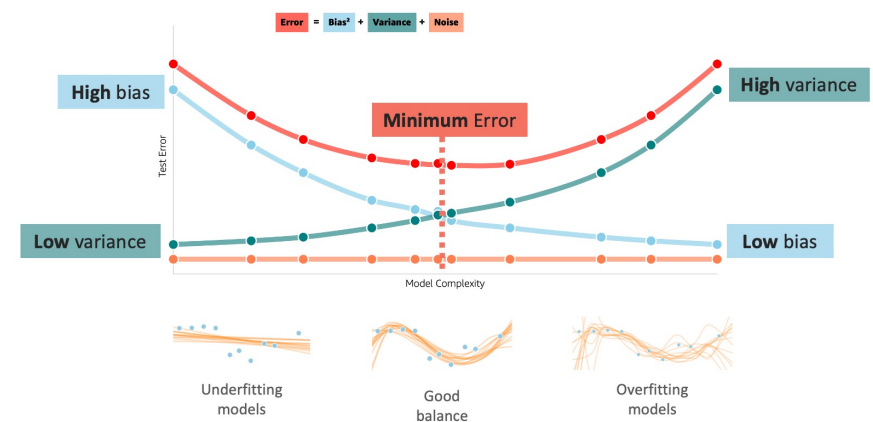


# Bias-Variance Tradeoff



# Bias-Variance Tradeoff

	Low Bias	High Bias
Low Variance	<p><b>Ideal:</b> accurate and consistent predictions (perfect fit)</p> 	<p><b>Underfitting:</b> predictions are far from the true value but clustered together</p> 
High Variance	<p><b>Overfitting:</b> some predictions are close, but highly scattered (model too sensitive)</p> 	<p><b>Worst case:</b> predictions are both inaccurate and inconsistent</p> 





# Bias-Variance Tradeoff


- Increasing bias reduce variance and vice-versa
- $Error = Bias^2 + Variance + Noise$
- The best model is where the error is reduced
- Compromise between bias and variance

# Linear regression

# Linear regression

Linear regression is a simple yet powerful method for predicting a numeric response from one or more independent variables.

- It is a **supervised algorithm** that learns to model a dependent variable,  $y$ , as a function of some independent variables (aka "**features**"),  $x_i$ , by finding a line (in 2D) or a hyperplane (in higher dimensions) that best "fits" the data.
- In general, we assume  $y$  to be some number and each  $x_i$  can be basically anything. For example:
  - Predicting the price of a house using the number of rooms in that house ( $y$  : price,  $x_1$ : number of rooms)
  - Predicting a person's weight based on their height and age ( $y$  : weight,  $x_1$ : height,  $x_2$ : age)



# Linear regression

In general, the equation for linear regression is:

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_p x_p + \varepsilon$$

Where,


$y$  = the dependent variable: the thing we are trying to predict

$x_i$  = the independent variable: the features our model uses to model  $y$ .

$\beta_i$  = the coefficients (aka “weights”) of our regression model.

These are the foundations of our model. They are what our model “learns” during optimization

$\varepsilon$  = the irreducible error in our model. A term that collects together all the unmodeled parts of our data




# Linear regression

Fitting a linear regression model involves finding the set of coefficients that best approximate  $y$  as a function of the input features.

- Although the true parameters of the data-generating process are unknown, we can estimate them from our data.
- Once we've estimated these coefficients ( $\hat{\beta}_i$ ), we can predict future values ( $\hat{y}$ ) using the equation:

$$\hat{y} = \hat{\beta}_0 + \hat{\beta}_1 x_1 + \hat{\beta}_2 x_2 + \cdots + \hat{\beta}_p x_p + \varepsilon$$

- Predicting future values (often referred as *inference*) is as simple as plugging the values of our features  $x_i$  into the model equation



# Linear regression


E.g., Let's fit a model to predict **housing price** (\$) using the size of the house **sqft** (in square-footage):

$$\hat{y} = \hat{\beta}_0 + \hat{\beta}_1 x_1$$

$$house_{price} = \hat{\beta}_0 + \hat{\beta}_1 * sqft$$

We'll start with a very simple model, predicting the price of each house to be just the average house price in our dataset, ~\$290,000, ignoring the different sizes of each house:


$$house_{price} = 290000 + 0 * 290000$$



# Linear regression


To **evaluate** our model's performance **quantitatively**, we plot the error for each observation. These errors, or **residuals**, represent the *distance between the actual observation and the predicted value* for that observation.

As shown in the plot, we can clearly see that our model exhibits a significant amount of error.



# Linear regression

- The goal of linear regression is to reduce prediction error by finding a line or surface that best fits our data.
- For a simple regression problem, this involves estimating the *intercept* ( $\hat{\beta}_0$ ) and *slope* ( $\hat{\beta}_1$ ) that define the relationship between our input features and the target variable.



There is still some error, but the general pattern is well captured. As a result, we can be reasonably confident that, by plugging in new values for square footage, our predicted house prices will be reasonably accurate.



# Linear regression

Predicting future values:

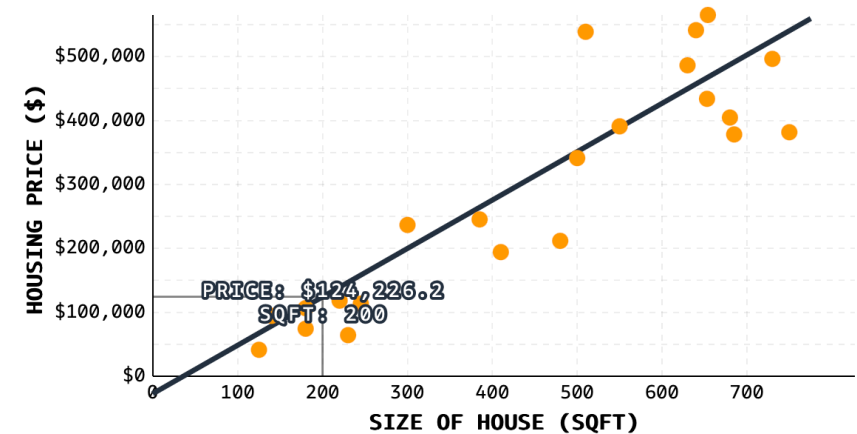
- We just plug in any  $x_i$  values into our equation!
- For our simple model, that means plugging in a value for **sqft** into our model:

**sqft** value = 200;

$$\hat{y} = -27153.8 + 756.9 * 200$$

$$\hat{y} = 124226$$

Our model predicts a house that is **200 square-feet** will  
**cost \$124,226**



# Linear regression

Predicting future values:


- We just plug in any  $x_i$  values into our equation!
- For our simple model, that means plugging in a value for **sqft** into our model:

**sqft** value = 580;

$$\hat{y} = -27153.8 + 756.9 * 580$$

$$\hat{y} = 411848$$

Our model predicts a house that is **580 square-feet** will  
**cost \$411,848**



# Model evaluation

They quantified how close the predicted value is to the true value. We'll fit our regression model to a set **training** data, and evaluate its performance on the **test** dataset.


Error	Formula	Range	Description
Total Absolute Error <b>TAE</b>	$TAE = \sum_{i=1}^n  Y_i - \hat{Y}_i $	$\geq 0$ (0 good)	<ul style="list-style-type: none"> <li>It is the sum of the absolute values of the difference between each true value and the predicted value (given by the regression line)</li> </ul>
Mean Absolute Error <b>MAE</b>	$MAE = \frac{1}{n} \sum_{i=1}^n  Y_i - \hat{Y}_i $	$\geq 0$ (0 good)	<ul style="list-style-type: none"> <li>It is the average of all absolute errors</li> <li>Easy to interpret: the units of MAE are the original of the data</li> <li><b>Not sensitive to outliers:</b> Treat larger and small errors equally</li> </ul>
Mean-Squared Error <b>MSE</b>	$MSE = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$	$\geq 0$ (0 good)	<ul style="list-style-type: none"> <li>It squares the distance between each true value and the predicted value (given by the regression line), summing the squared values, and then dividing by the number of data points</li> <li><b>Sensitive to outliers:</b> It punishes large errors more</li> </ul>

$n$  = number of data records

$Y_i$  = true values

$\hat{Y}_i$  = predicted values

$Y$  = mean value of true values;  $Y = \frac{1}{n} \sum_{i=1}^n Y_i$



# Model evaluation

They quantified how close the predicted value is to the true value. We'll fit our regression model to a set **training** data, and evaluate its performance on the **test** dataset.


Error	Formula	Range	Description
Root Mean-Squared Error RMSE	$RMSE = \sqrt{\frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2}$	$\geq 0$ (0 good)	<ul style="list-style-type: none"> <li>It is the root of the Mean-Squared Error (MSE)</li> <li>Easy to interpret: the units of RMSE are the original of the data</li> <li>Sensitive to outliers: It punishes large errors more</li> </ul>
Root Squared Error $R^2$	$R^2 = 1 - \frac{\sum_{i=1}^n (Y_i - \hat{Y}_i)^2}{\sum_{i=1}^n (Y_i - \bar{Y})^2}$	$0 \leq R^2 \leq 1$ (0 bad, 1 good)	<ul style="list-style-type: none"> <li>It summarizes how well a model fits the data.</li> <li>It represents the % of the variance in <math>y</math> explained by the features <math>x</math>. It takes the MSE (numerator) and the variance of the mean (denominator)</li> <li><math>R^2 = 1</math>: model capture 100% of the variance</li> <li><math>R^2</math> Negative value: the predicted model performs worse than a simple average of the original data</li> </ul>

$n$  = number of data records

$Y_i$  = true values

$\hat{Y}_i$  = predicted values

$\bar{Y}$  = mean value of true values;  $\bar{Y} = \frac{1}{n} \sum_{i=1}^n Y_i$



# Model evaluation

$$\hat{y} = \hat{\beta}_0 + \hat{\beta}_1 x$$


$$\hat{\beta}_0 = 6$$

$$\hat{\beta}_1 = 0.01$$

$$\hat{y} = 6 + 0.01x$$

$$MSE = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2 = 5.22$$

$$R^2 = 1 - \frac{\sum_{i=1}^n (Y_i - \hat{Y}_i)^2}{\sum_{i=1}^n (Y_i - \bar{Y})^2} = 0.01 = 1\%$$



# Model evaluation

$$\hat{y} = \hat{\beta}_0 + \hat{\beta}_1 x$$

$$\hat{\beta}_0 = 2$$

$$\hat{\beta}_1 = 0.76$$

$$\hat{y} = 2 + 0.76x$$

$$MSE = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2 = \mathbf{0.71}$$

$$R^2 = 1 - \frac{\sum_{i=1}^n (Y_i - \hat{Y}_i)^2}{\sum_{i=1}^n (Y_i - \bar{Y})^2} = \mathbf{0.87 = 87\%}$$

